

MorphSeq Param Sweep with Splines; Dec-17-2024

javascript

type: daily

date: 2024-12-17

project: #morphSeq

date: Dec-17-2024

Dec 17 2024

MorphSeq Param Sweep with Splines

Previous Parameter Sweep (where we left off)

This is following up on the previous parameter sweep. Since then I have generated splines and different ways to measure colinearity WITHIN and across the splines for perturbations.

Main takeaways from previous param sweep:

Main Takeaways

- Classification is not a trivial task and there is a spread in model performance
- We can have "pretty-morphology" spaces that are good at classification
- Right now we don't have a very good way of measuring "pretty"
- NT-Xent this was the top performing model.
- We don't have a good way to measure "prettiness" nor different type of it
 - There is some signal because Clustering coefficient (measure of cliqueness) did have a trend potentially.
- Are there model architectural decisions that would impact classification reproducibility

Questions:

- why do all the "prettiest" best models have the highest value for Beta

- I eventually how well deep learning vs log reg for classification
 - Does Aesthetic of the model matter
 - Stability
 - Should we measure these dim reduction space versus
-

Main Takeaway

Key Players

I used splines using [@einbeckLocalPrincipalCurves20E05](#) to fit splines to the data we wanted to Measure

- Tree Likeness (Prettiness), done by looking at dispersion (distance to centroid) of splines
 - dispersion_first_n and dispersion_last_n is how close the points are in the first n points along the datasets splines (or lastn)
 - Note i used n=5
 - disp_coefficient is a fitting a linear model to the dispersion ALONG the spline
 - This is to make sure that splines (and therefore points) are spreading out over time
 - SegmentColiniarity_mean_within (all or hld): point moving in same direction
 - This is looking WITHIN a dataset and looking at pairwise between perts how these points move in the same direction (cosine similiarty)
- Robustness
 - SegColinearity, is looking at a model INSTANCE (each model has all or hld (holding out lm1b and gdf3)) and then looking at how for the same pert they move in the same direction.
 - Aligned MSE using the quaternion implementaion of Kabsch aligning algorithm we look at how the MSE between like perturbation for a model instance is.
 - This is very stringent but a good measure.
- Initial MSE
 - How without aligning the MSE
 - used as an internal check to make sure Alignmen tis working AND to see if alignment is actually necessary.
- Classification:
 - F1 Score a measure of how good model is at classification
 - for more info see previous param sweep (referenced above)

Conclusion

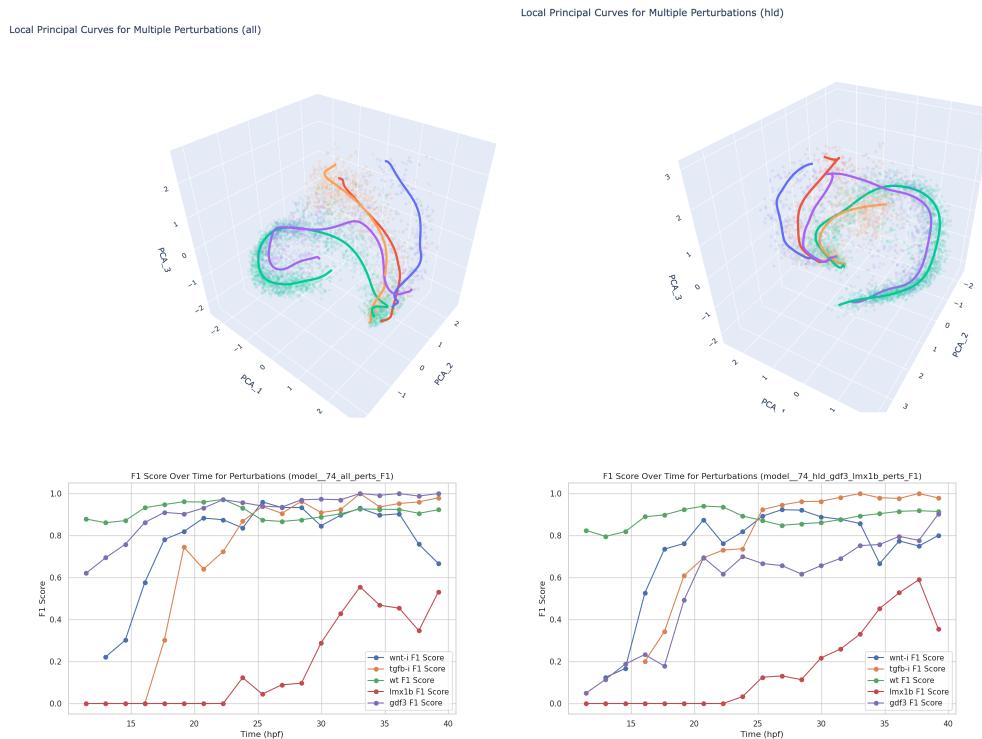
Chosen Model (74):

I think Model 74 (or another model in its class) should be chosen for these reasons)

Summary of why model 74

- I choose model 74 because it still has tree-like structure, is stable, and it has good classification accuracy
- Reading below we see that it does a better job extracting data than other models e.g. model 97 (not necessarily those in its same tanche e.g. 77)
- It even does better than the previous model we were using!!!

Sel Model_Index 74

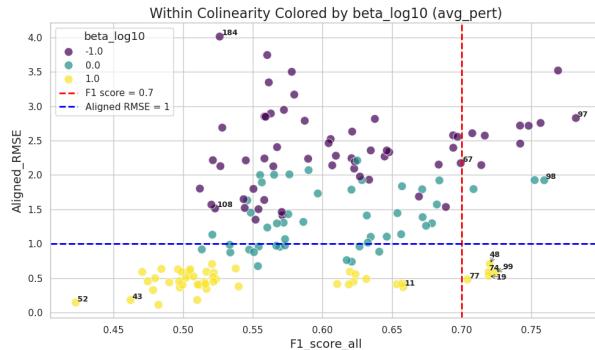


Tradeoff of Prettiness vs Classification

- There is a tradeoff on prettiness vs classification accuracy, though this is slight. It is MOST definitely worth the small decrease in classification accuracy to get a more interpretable morphology space.
 - See [MorphSeq Param Sweep with Splines; Dec-17-2024 > Tradeoff of Prettiness vs Classification](#) for more detail

Beta Drives Prettiness AND Robustness

- It can be seen that the Beta variable drives prettiness and robustness
- Higher values of the beta variable makes the model make the individual dimensions in the embedding space more iid (independent)
 - It also clearly has the property of making morphology space prettier
 - potentially by smoothing things out.



- This is a surprising result as prettiness/tree likeness is NOT guaranteed to be correlated with robustness/stability.
- High beta models are so robust that they actually don't need to be aligned! (though you should still align them lol)

New Spline Measurement do capture things we wanted them too

- Implementing splines allowed us to finally get at what was the drivers of prettiness
 - previously not able to

Concerns

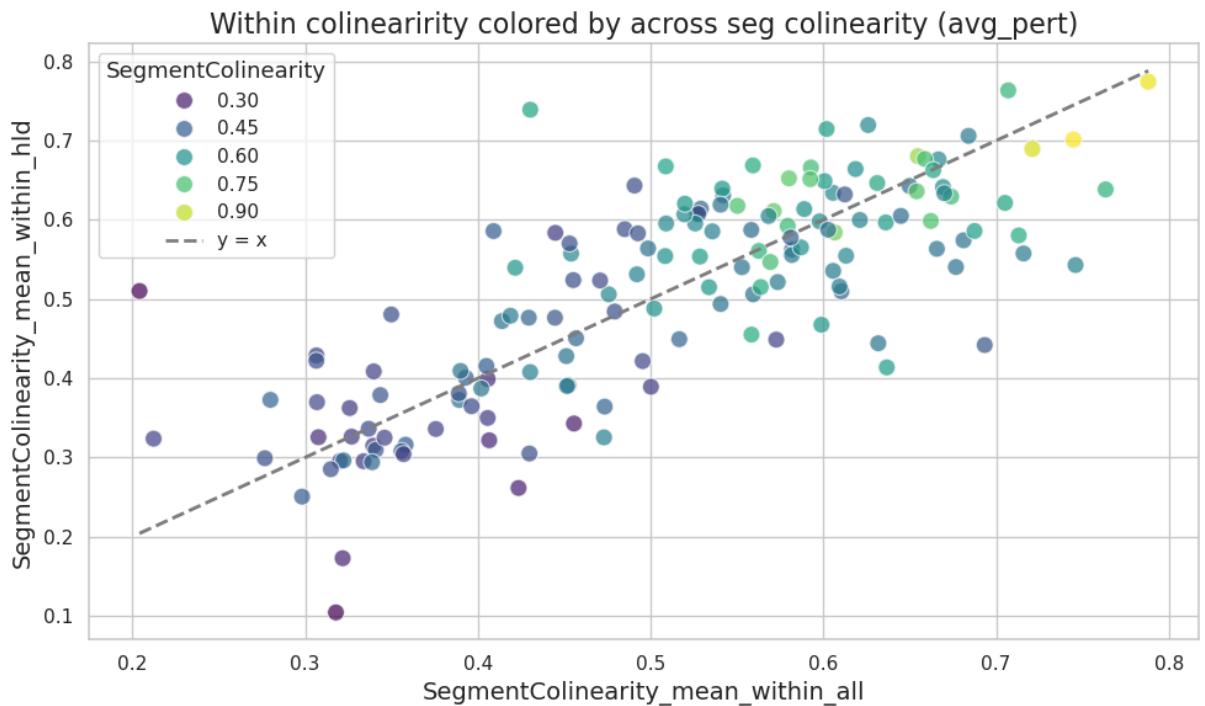
- sometimes splines don't work, but this is when they aren't pretty (tight turns, not tree shaped) see model 98
- but these might be inflating the Aligned MSE values but not really worried about this.
 - because when they fail it's because data isn't tree like.

Next Steps

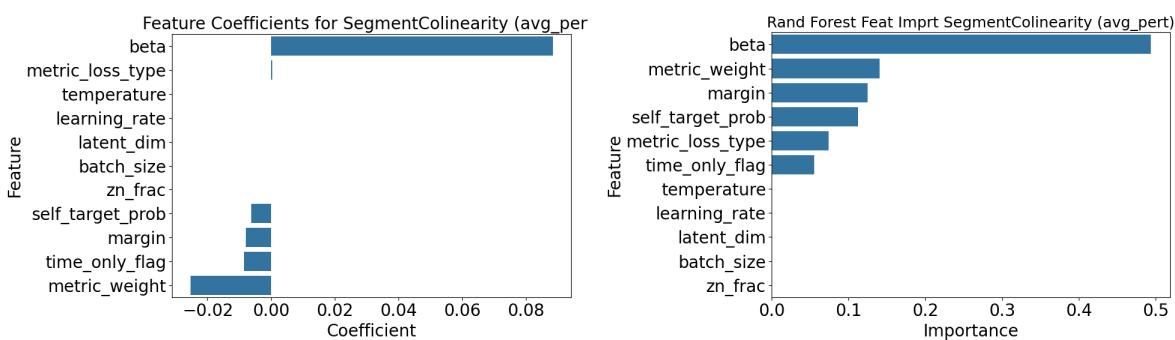
- Consider other measurements that don't need splines (none come to mind)
- I see no obvious change, but if we were to make the dispersion_metric first_n and last_n could be calculated from one linear equation, the y intercept could be first_n and the multiple the slope by the last index (which is the dispersion coefficient here actually) and then add the slope to get last_n
- Of course we need to find a way of projecting points onto the spline and then time across it
 - measure variance across spline
 - measure variability of time across spline

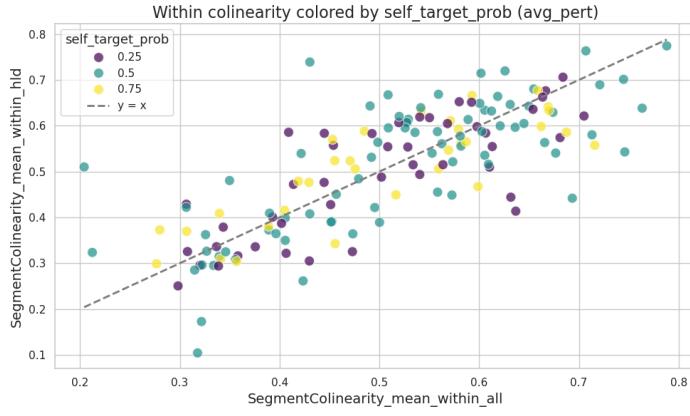
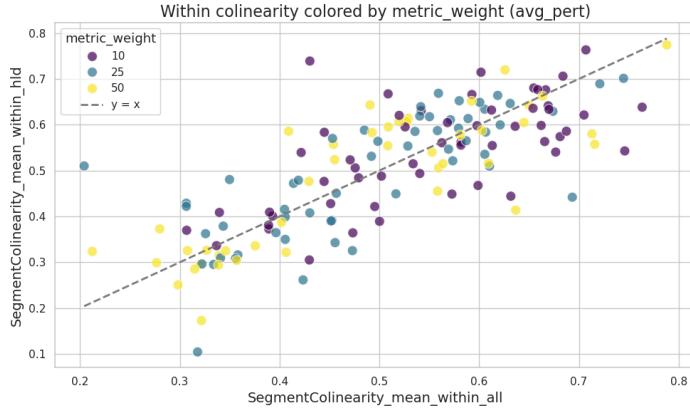
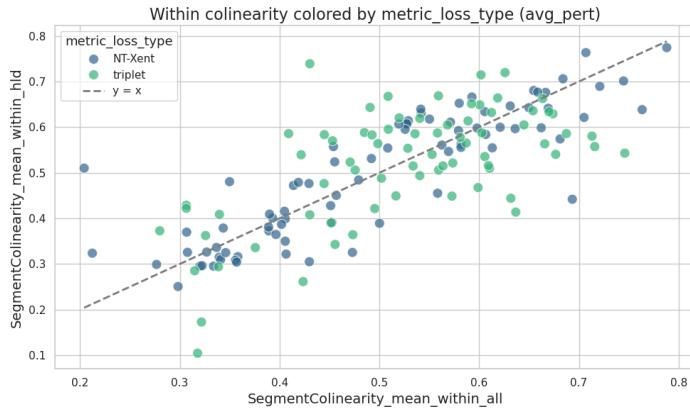
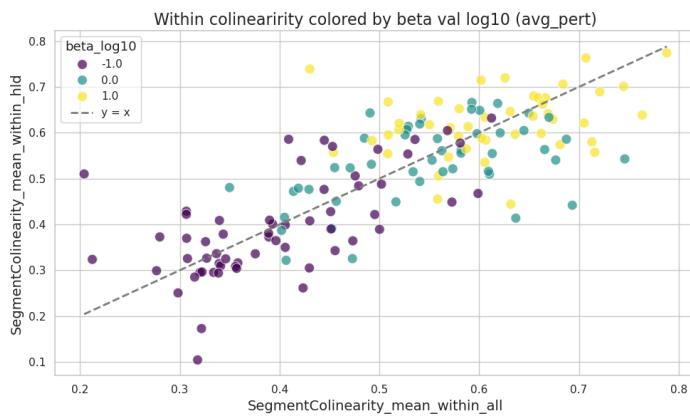
Collinearity

- Here we see that if a model is not colinear within a dataset it wont be in the other,
- Furthermore, that if it has a high collinearity within one then it will be colinear with the held out data! (this can be seen by coloring the points by ACROSS dataset collinearity). This is the first actual evaluation of
- this means that if you project into a space that is already linear, you will get linear manifolds.



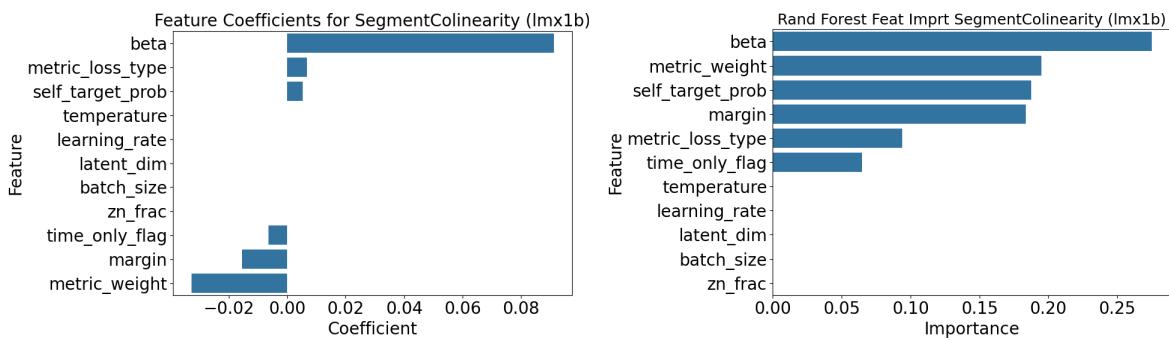
beta as causative for collinearity AND Robustness



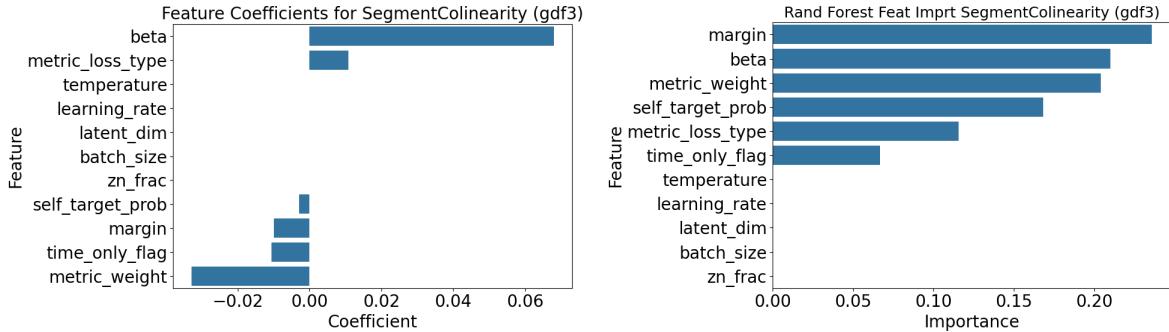


- note that NT-Xent has tighter distributions around $y=x$ (more stable)

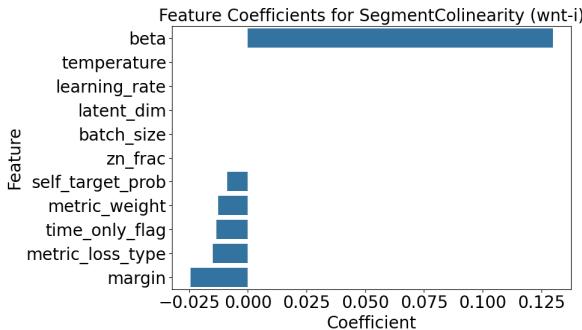
- lm1b



- gdf3



![[Projects/morphSeq/Journal/Dec-2024/attachments/MorphSeq Param Sweep with Splines;Dec-17-2024/image 15.png|300]][Projects/morphSeq/Journal/Dec-2024/attachments/MorphSeq Param Sweep with Splines;Dec-17-2024/image 16.png|300]]

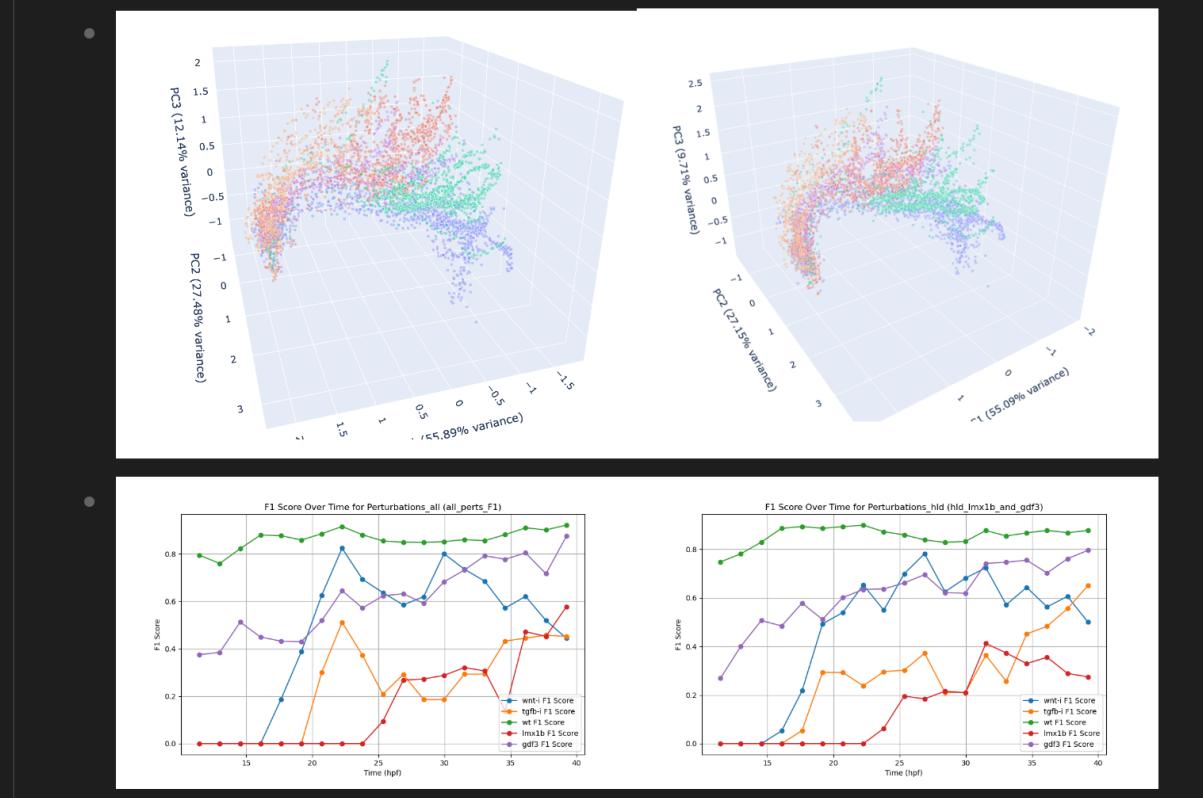


MSE Initial -> Aligned

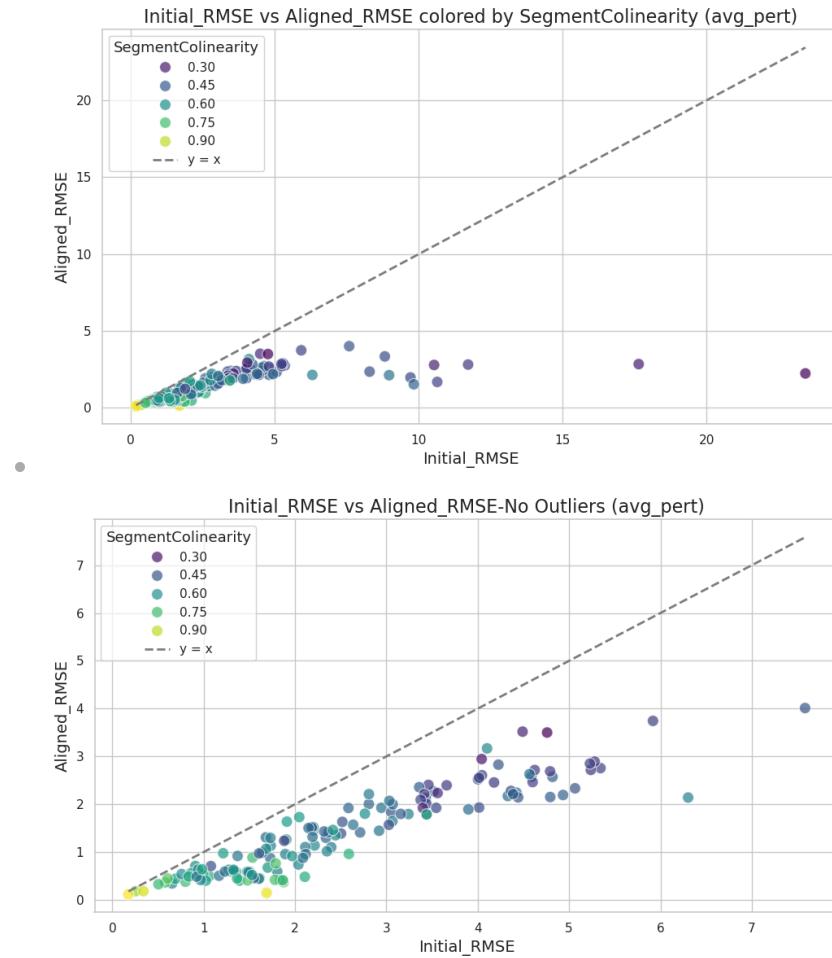
- Here I Noticed that for one of the models that did the best in terms of colinearity metrics, that the PCA was RIGHT on top of eachother pre alignemnt which was not garunteed at all. (this would imply that PCA is generating the same eigenvectors acorss training regimes)
- example 71 is a prime example!
 - [MorphSeq Parameter Sweep Takeaways;Nov-28-2024](#)

• 71: F1 0.72

- time only was used

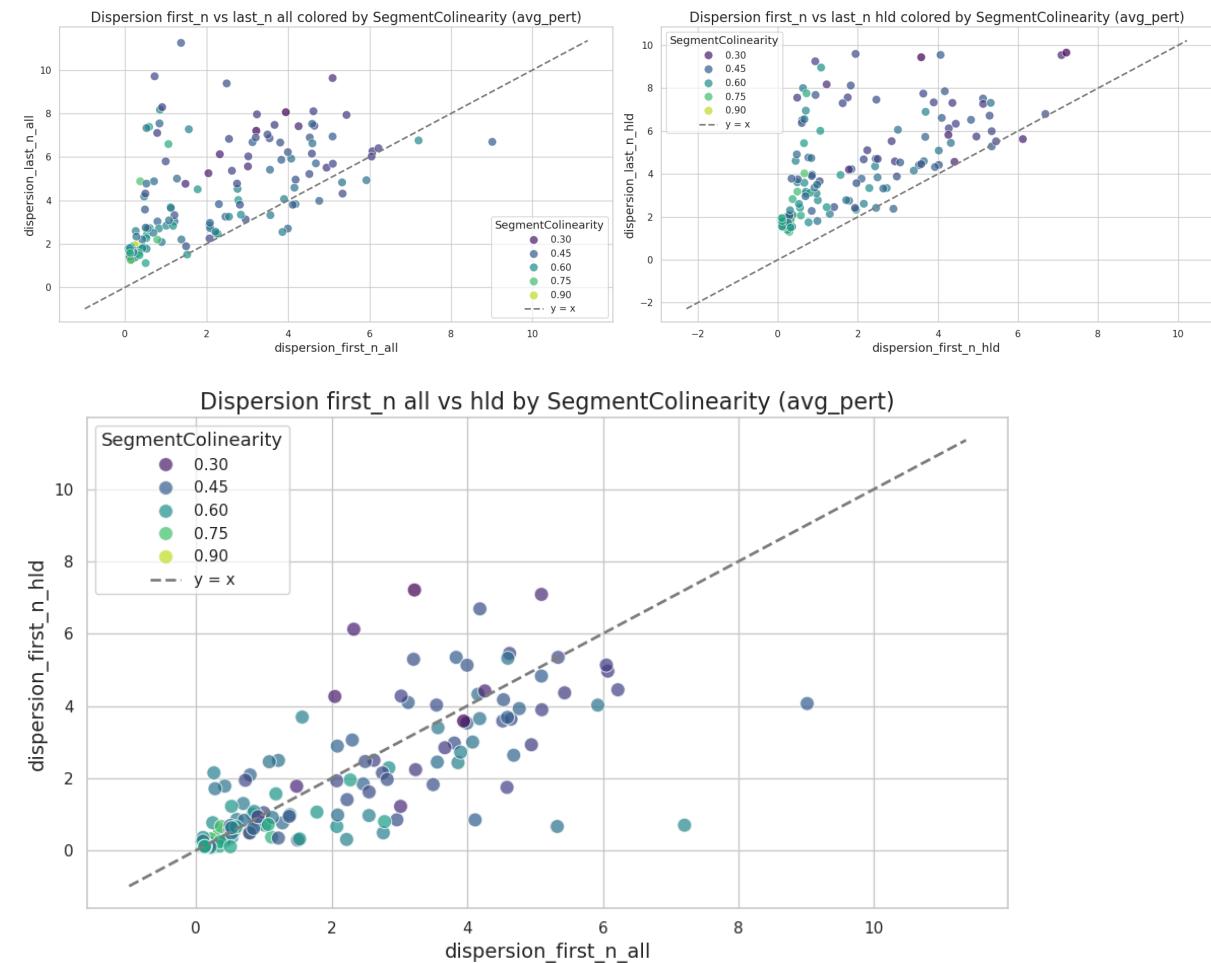


- This is confirmed here! that yes if the model is colinear then it will already be aligned!
 - Note that all the lines are under the $y=x$ line which means Initial RMSE < Aligned RMSE
 - (second plot is removing outliers)

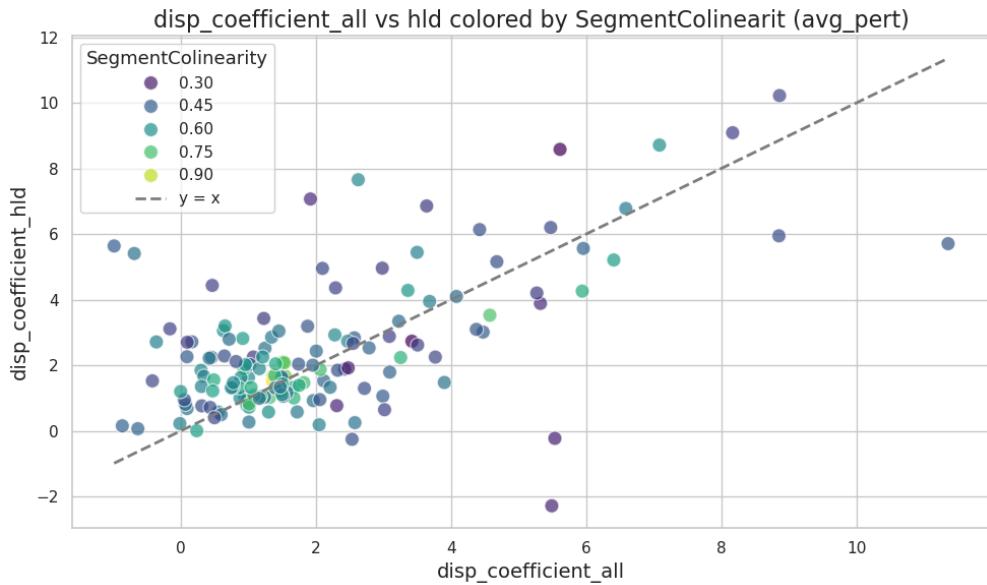


Dispersion (tree likeness)

- we wanted to see how tree like our models by looking at the first n and last n points and seeing how spread apart they are, we want first n to be close to zero and last n to be bigger. Additionally we use dispersion coefficient which simply looks at dispersion along the spline. higher dispersion_coef means higher spread across the spline. (note n = 5)
- main takeaway from these plots is that as long as the model is average or better at the segment collinearity metrics then it will be tree and that generally dispersion increases over the spline.
- here we see that for models that do ok with linearity metrics have a first_n close to 0 AND last n around 2
 - also note that when held out model the beginning points (first_n) become more spread out as well

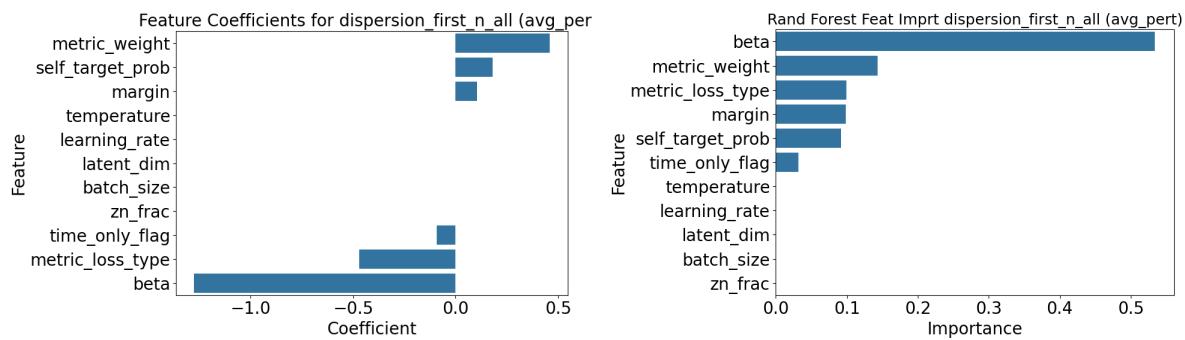


- We can see that model becomes slightly more disperse (higher dispersion coefficient in hld) when holding out perturbations.



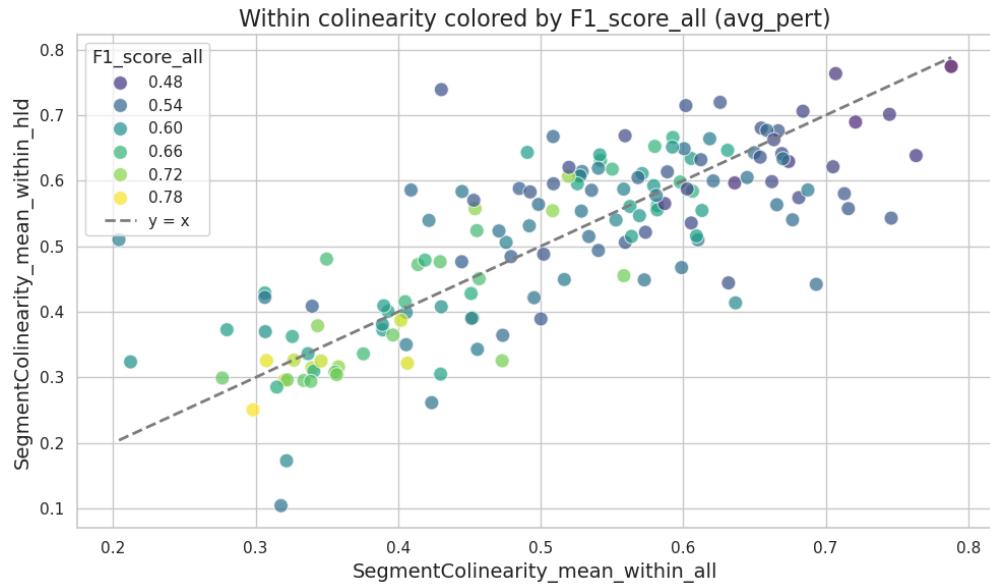
drivers of Dispersion:

- Higher beta --> more linear --> closer first_n --> more tree like

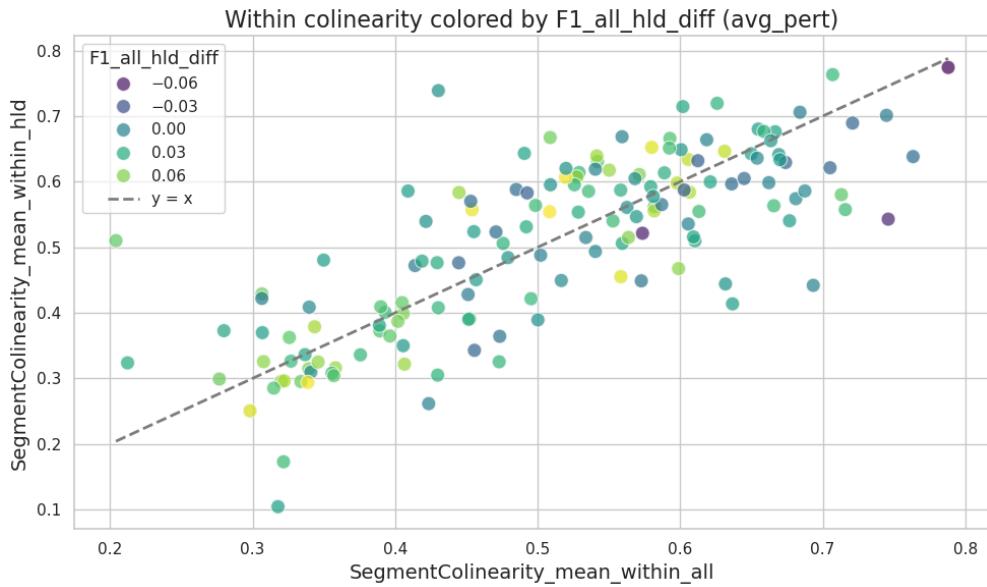


Prettiness vs Classification Accuracy

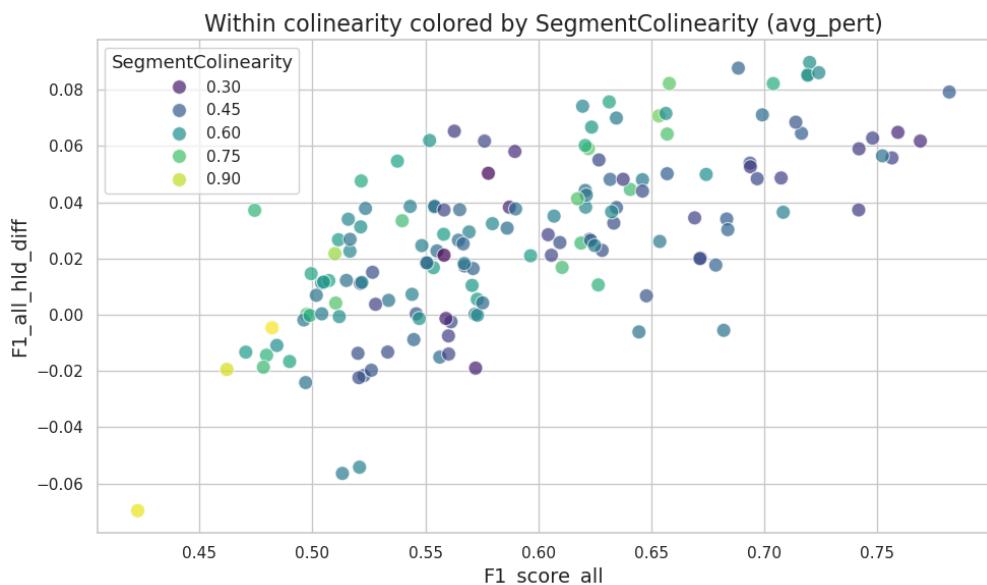
- clear trend that:
 1. classification can come at the cost of prettiness and vice versa
 2. for a given f1 score more colinearity decreases classification robustness
- coloring segcoliniarity by F1 score showes that there is generally a tradeoff that lower F1 scores (bluer) corresponds to higher colinearity metrics



- Hard to draw a trend between F1 score difference (robustness) and seg colinearity

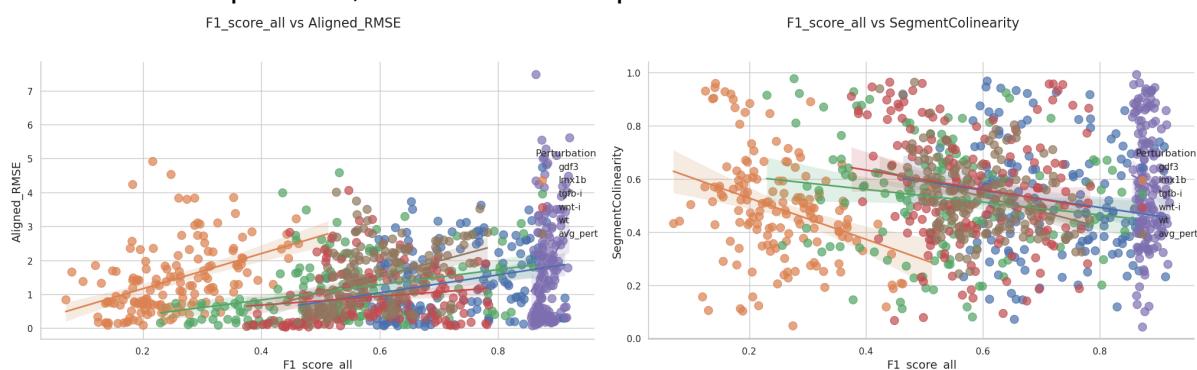


- Slight signal but you can SEE that for a given F1 score higher colinearity means that there will be more of a performance loss wrt. F1 score robustness...
 - though the tradeoff isn't anything crazy...



Classification on per Pert level and Prettiness

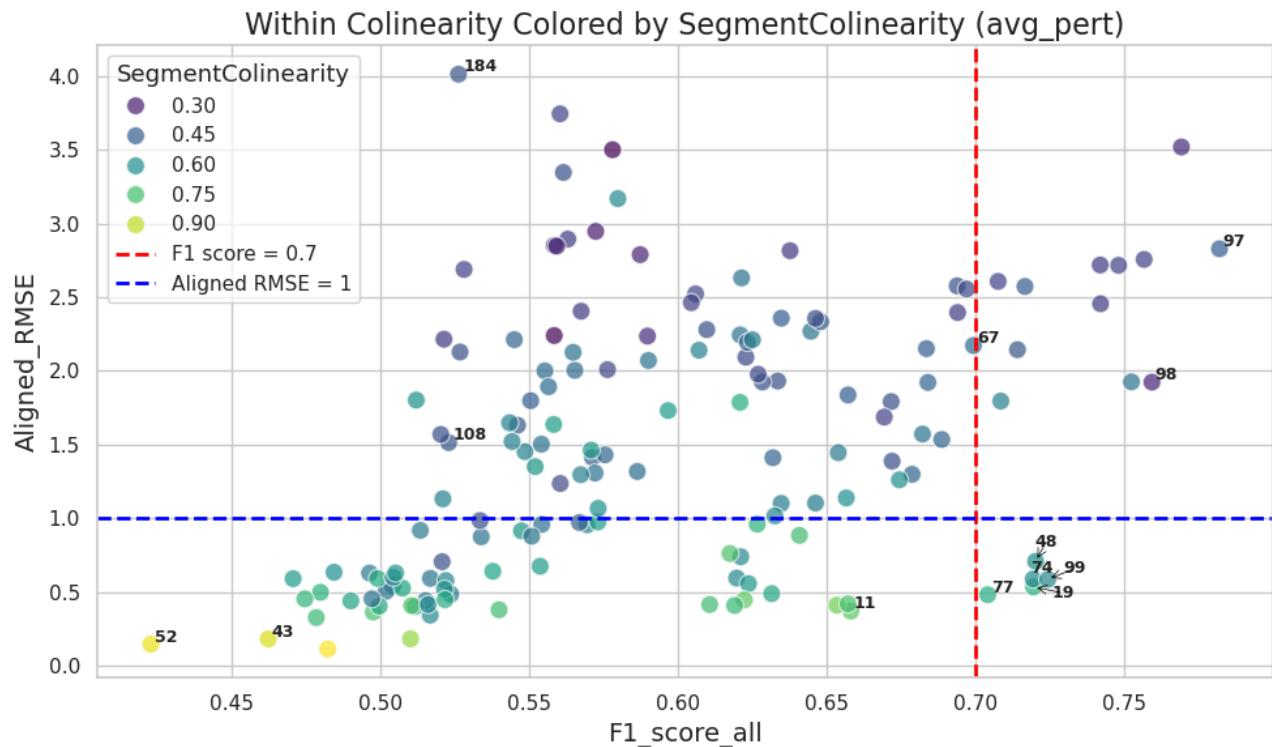
- Trends are recapitulated, tradeoff between prettiness and classification



Picking models

First pass with metrics we care about F1 and Align MSE (prettiness)

- I think having **F1 score > .7** and **Aligned MSE < 1** are good acceptable models (look at blue and red lines)
- We care about F1 score (classification), Aligned RMSE (robustness AND prettiness), I am coloring by other things too, to make sure they have good properties.
- Note we know Aligned RMSE is associated with robustness AND prettiness) because of [MorphSeq Param Sweep with Splines; Dec-17-2024 > MSE Initial -> Aligned](#)



Sel Models from different parts of the plot used to select models:

using Aligned RMSE and F1 score as discriminators

- Best models (98,48,19,74,77)

Aligned MSE < 1 F1 score >.7 (BEST MODELS)																
# self_df = merged_df_avg[(merged_df_avg["Aligned_RMSE"] <= 1) & (merged_df_avg["F1_score_all"] > 0.7)]																
✓	0.05	model_index	Perturbation	metric_loss_type	beta	time_only_flag	self_target_prob	margin	metric_weight	F1_score_all	F1_score_hid	Aligned_RMSE	SegmentColinearity	SegmentColinearity_mean_within_all	dispersion_first_n_all	disp_coefficient_all
479	99	avg_pert	NT-Kent	10.0	0	0.25	5.0	50	0.724121	0.638041	0.585608	0.567616	0.519465	0.425045	1.162380	
245	48	avg_pert	NT-Kent	10.0	0	0.25	1.0	50	0.720055	0.630332	0.707545	0.610988	0.508386	0.456199	0.890601	
101	19	avg_pert	NT-Kent	10.0	0	0.25	0.1	50	0.719329	0.633751	0.534260	0.710807	0.526650	0.451727	0.934226	
365	74	avg_pert	NT-Kent	10.0	0	0.25	2.0	50	0.719221	0.634051	0.587776	0.587008	0.453926	0.530416	0.300422	
377	77	avg_pert	NT-Kent	10.0	0	0.50	2.0	50	0.703967	0.627768	0.481163	0.6597584	0.558336	0.347210	1.751113	

- Second Tranche (97,98)

Aligned MSE >1 F1 score >.7 (Subset)																
Python																
merged_df_avg[merged_df_avg["Aligned_RMSE"] > 1] & (merged_df_avg["F1_score_all"] > 0.7) & (merged_df_avg["model_index"].isin([97,98])) (inp_cols).sort_values(by="F1_score_all", ascending=False)	0.05															
model_index	Perturbation	metric_loss_type	beta	time_only_flag	self_target_prob	margin	metric_weight	F1_score_all	F1_score_hid	Aligned_RMSE	SegmentColinearity	SegmentColinearity_mean_within_all	dispersion_first_n_all	disp_coefficient_all		
467	97	avg_pert	NT-Xent	0.1	0	0.25	5.0	50	0.781858	0.702657	2.827774	0.47553	0.298167	9.013367	-0.986660	
473	98	avg_pert	NT-Xent	1.0	0	0.25	5.0	50	0.759139	0.694277	1.923200	0.325243	0.406500	1.483292	2.308945	

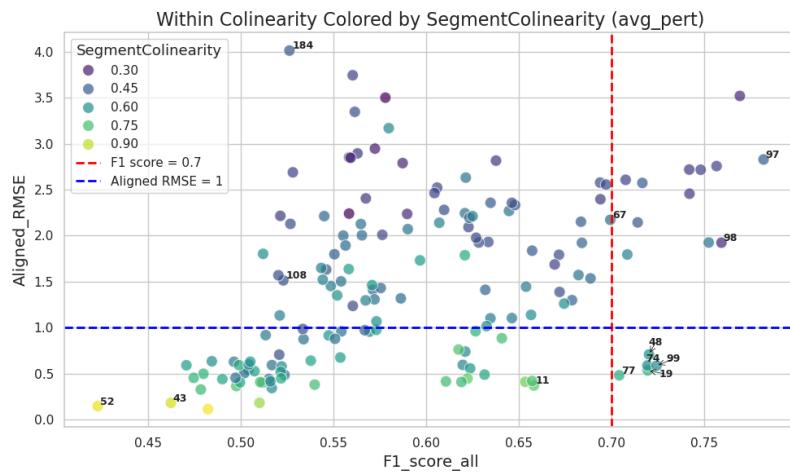
- Third Tranche (11,43,52)

Aligned MSE <=1 F1 score <.7 (Subset)																
Python																
merged_df_avg[(merged_df_avg["Aligned_RMSE"] <= 1) & (merged_df_avg["F1_score_all"] < 0.7) & (merged_df_avg["model_index"].isin([52,43,11])) (inp_cols).sort_values(by="F1_score_all", ascending=False)]	0.05															
model_index	Perturbation	metric_loss_type	beta	time_only_flag	self_target_prob	margin	metric_weight	F1_score_all	F1_score_hid	Aligned_RMSE	SegmentColinearity	SegmentColinearity_mean_within_all	dispersion_first_n_all	disp_coefficient_all		
59	11	avg_pert	NT-Xent	10.0	0	0.25	0.1	25	0.657936	0.576697	0.369914	0.779365	0.580015	0.146782	1.528816	
215	43	avg_pert	NT-Xent	10.0	1	0.50	1.0	25	0.462121	0.481579	0.180996	0.920061	0.720897	0.128076	1.363076	
257	52	avg_pert	NT-Xent	10.0	1	0.50	1.0	50	0.422659	0.492315	0.146000	0.940644	0.788008	0.266395	1.485208	

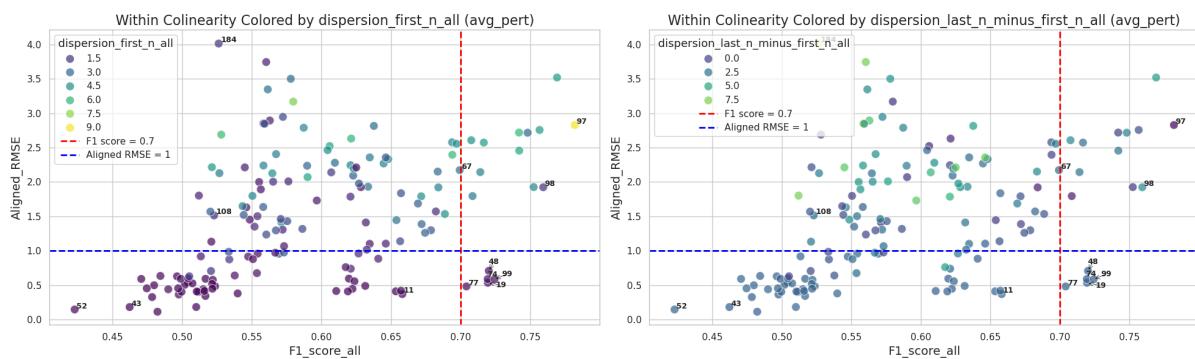
- Fourth Tranche (67, 184, 108)

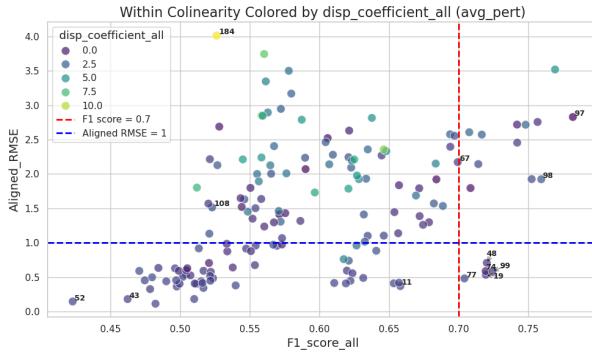
Aligned MSE >1 F1 score <.7 (Subset)																
Python																
merged_df_avg[(merged_df_avg["Aligned_RMSE"] > 1) & (merged_df_avg["F1_score_all"] < 0.7) & (merged_df_avg["model_index"].isin([108,67,184])) (inp_cols).sort_values(by="F1_score_all", ascending=False)]	0.05															
model_index	Perturbation	metric_loss_type	beta	time_only_flag	self_target_prob	margin	metric_weight	F1_score_all	F1_score_hid	Aligned_RMSE	SegmentColinearity	SegmentColinearity_mean_within_all	dispersion_first_n_all	disp_coefficient_all		
323	67	avg_pert	NT-Xent	0.1	0	0.50	2.0	25	0.699149	0.499992	0.321609	3.577945	1.022507			
881	184	avg_pert	triplet	0.1	0	0.25	5.0	10	0.526068	0.628076	4.012054	0.443271	0.444577	1.374297	11.356005	
515	108	avg_pert	triplet	0.1	0	0.50	0.1	10	0.522775	0.544393	1.513283	0.435893	0.612368	1.213509	2.058222	

- assuring models have nice segment colinearity

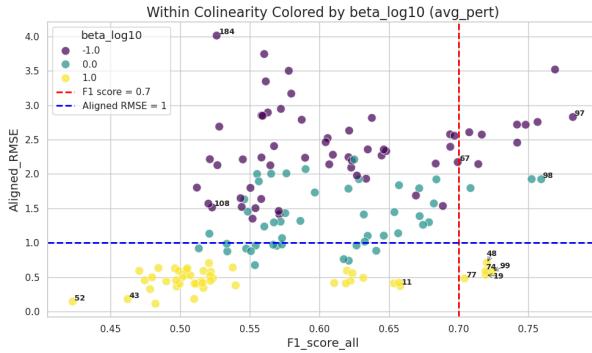


- assuring models have tree like structure low disp_first_n and non abnormal dispersion_coefficient





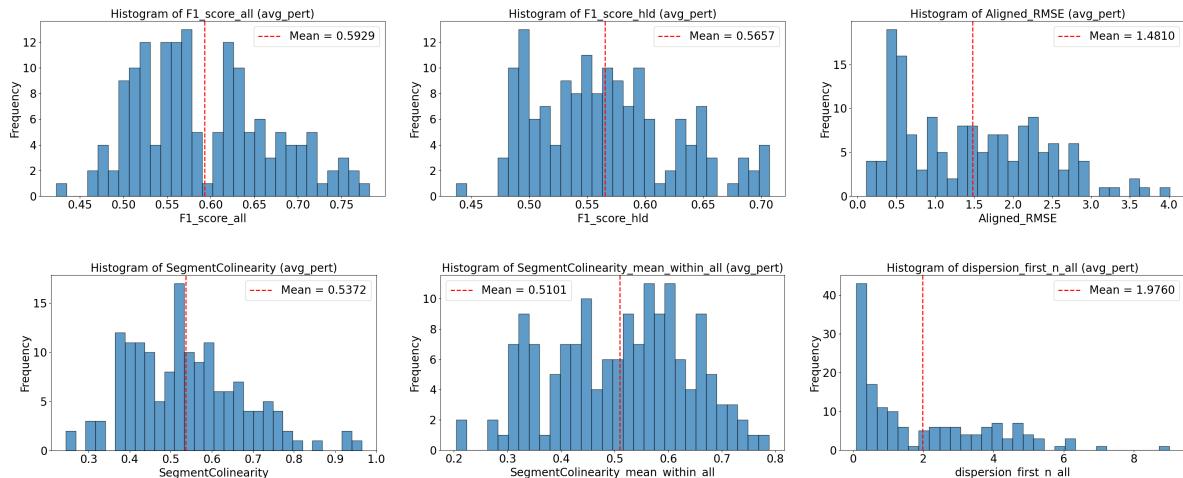
- we can see beta is driving the prettiness



Extracting the top model models (98,48,19,74,77):

Aligned MSE < 1 F1 score >.7 (BEST MODELS)															
# self_df = merged_df_avg[(merged_df_avg["Aligned_RMSE"] <= 1) & (merged_df_avg["F1_score_all"] > 0.7)]	Python														
# self_df[imp_cols]	merged_df_avg[(merged_RMSE <= 1) & (merged_df_avg["F1_score_all"] > 0.7)][imp_cols].sort_values(by="F1_score_all", ascending=False)														
0.08															
model_index	Perturbation	metric_loss_type	beta	time_only_flag	self_target_prob	margin	metric_weight	F1_score_all	F1_score_hld	Aligned_RMSE	SegmentColinearity	SegmentColinearity_mean_within_all	dispersion_first_n_all	disp_coefficient_all	
479	99	avg_pert	NT-Xent	10.0	0	0.25	5.0	0.724121	0.638041	0.585608	0.587616	0.519465	0.425045	1.162380	
245	48	avg_pert	NT-Xent	10.0	0	0.25	1.0	50	0.720055	0.630332	0.707545	0.610888	0.508386	0.456199	0.890601
101	19	avg_pert	NT-Xent	10.0	0	0.25	0.1	50	0.719329	0.633751	0.534260	0.710807	0.526650	0.451727	0.934226
365	74	avg_pert	NT-Xent	10.0	0	0.25	2.0	50	0.719221	0.634051	0.587776	0.587008	0.455926	0.530416	0.300422
377	77	avg_pert	NT-Xent	10.0	0	0.50	2.0	50	0.703967	0.621768	0.481153	0.697584	0.558336	0.347210	1.751113

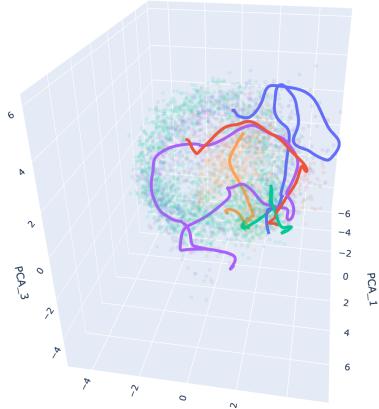
- None of them use the time only flag, and ALL of them are NT-Xent
- we can see that model is in top quartile for Classification metrics and has very close first_n, and is average in terms of segment colinearity



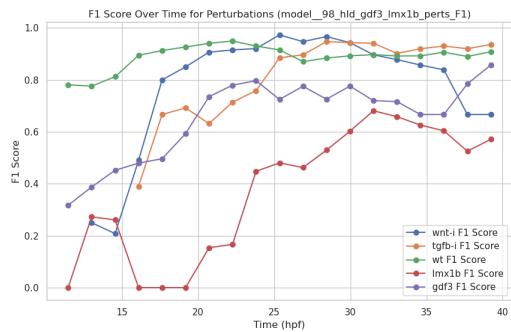
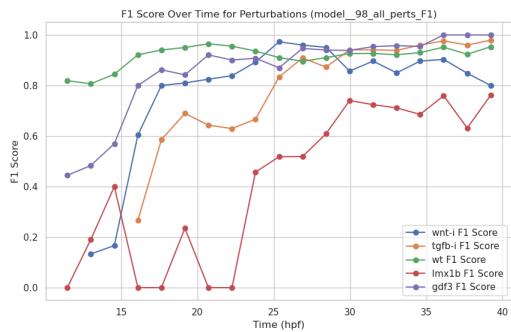
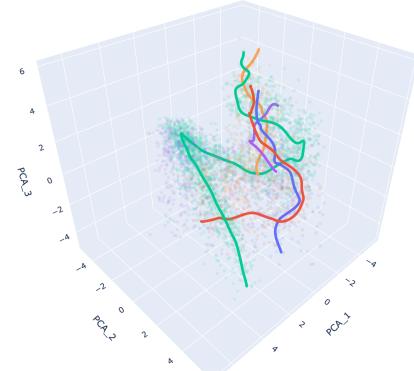
Looking at Best Models (Lower right of plot)

Sel Model_Index 98:

Local Principal Curves for Multiple Perturbations (hld)

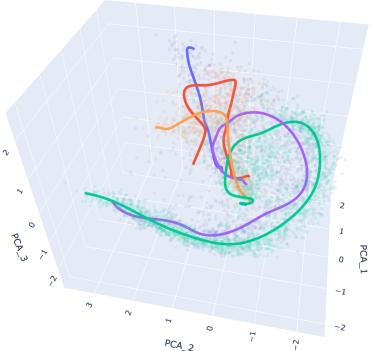


Local Principal Curves for Multiple Perturbations (all)

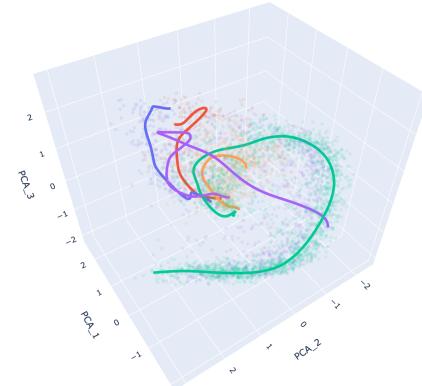


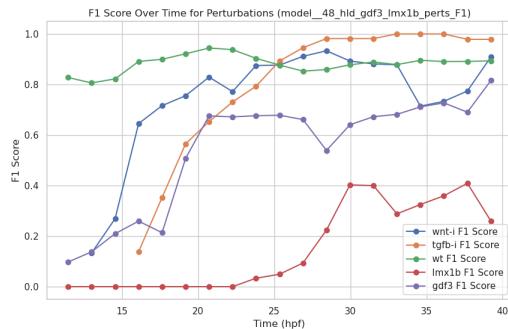
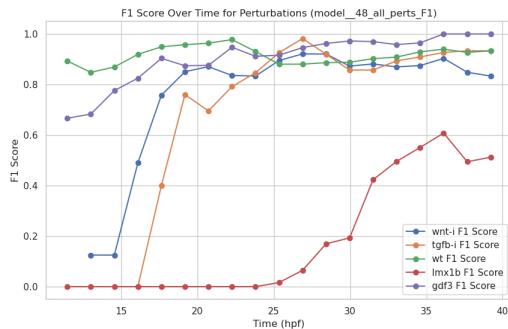
Sel Model_Index 48

Local Principal Curves for Multiple Perturbations (all)



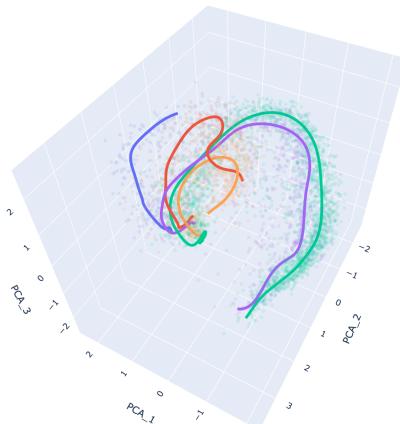
Local Principal Curves for Multiple Perturbations (hld)



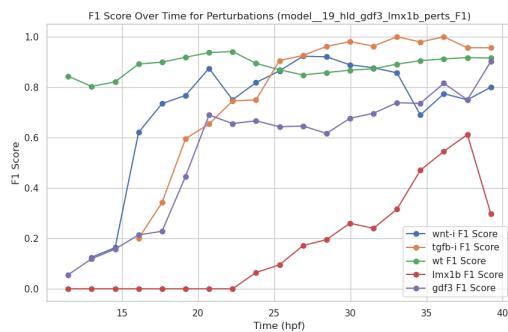
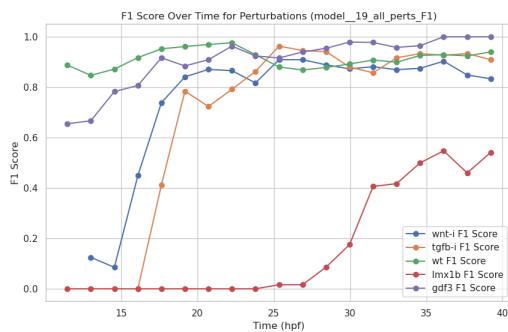
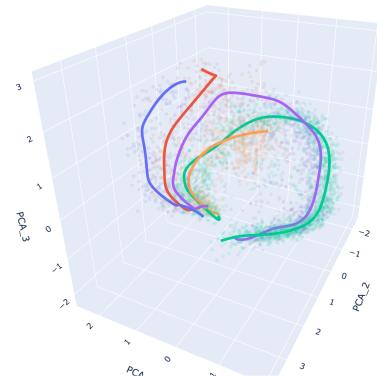


Sel Model_Index 19

Local Principal Curves for Multiple Perturbations (all)



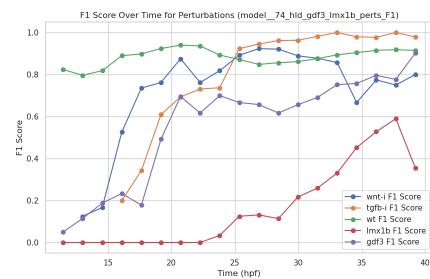
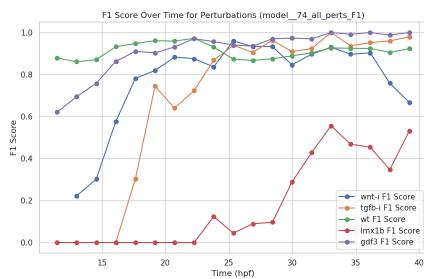
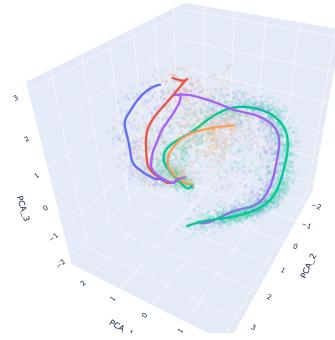
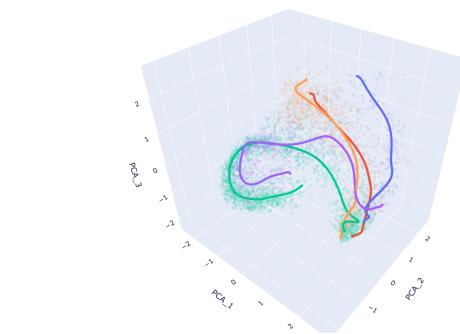
Local Principal Curves for Multiple Perturbations (hld)



Sel Model_Index 74

Local Principal Curves for Multiple Perturbations (all)

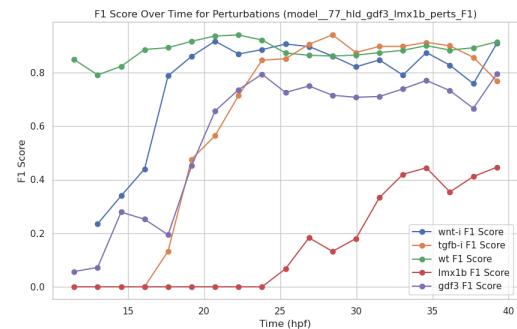
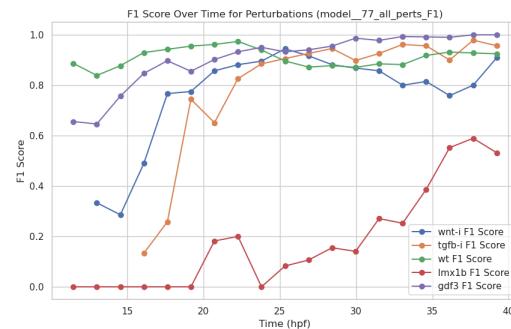
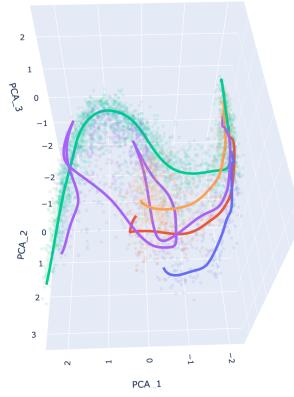
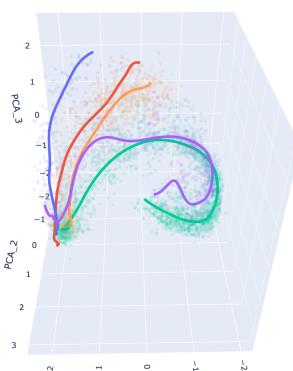
Local Principal Curves for Multiple Perturbations (hld)



Sel Model_Index 77

Local Principal Curves for Multiple Perturbations (all)

Local Principal Curves for Multiple Perturbations (hld)

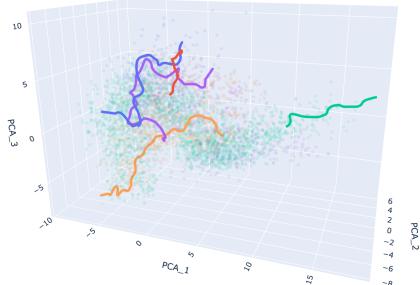


Looking at Second Traunche (97,98; Top Right)

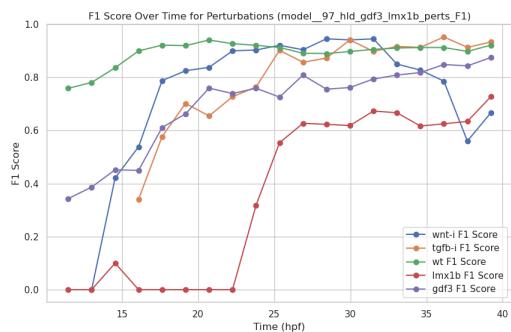
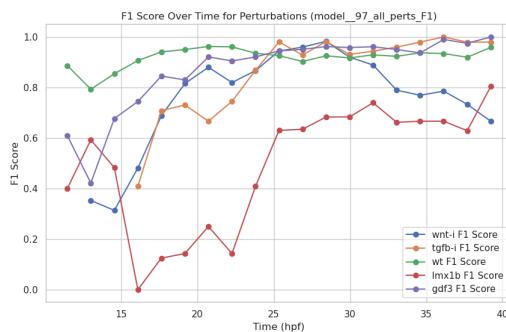
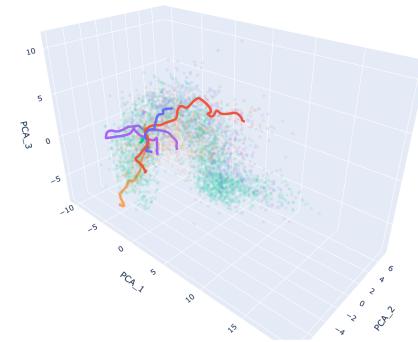
- These models have the highest F1 score 97 has .78 which is the higher

Sel Model_Index 97

Local Principal Curves for Multiple Perturbations (all)

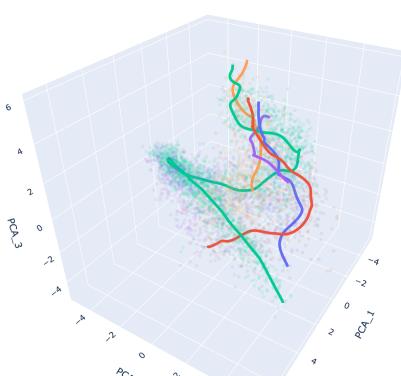


Local Principal Curves for Multiple Perturbations (hld)

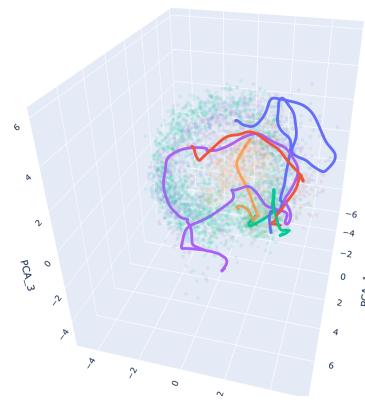


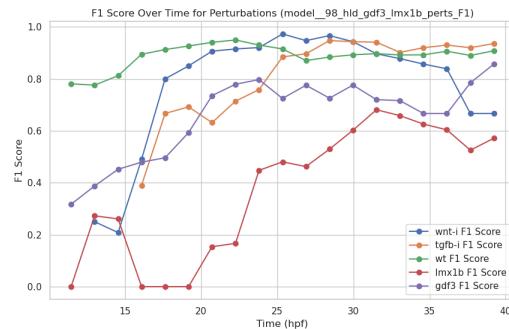
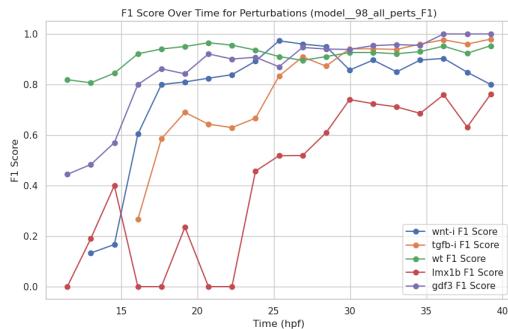
Sel Model_Index 98

Local Principal Curves for Multiple Perturbations (all)



Local Principal Curves for Multiple Perturbations (hld)

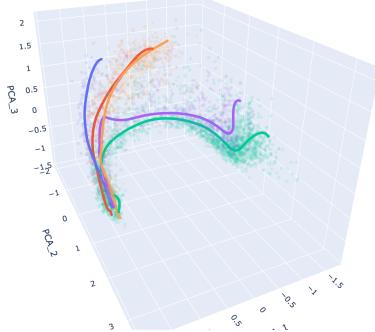




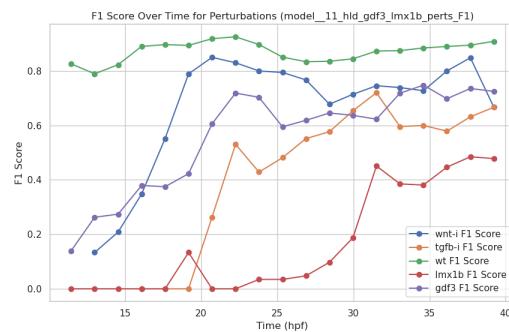
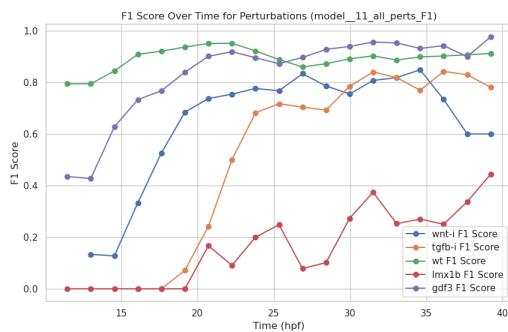
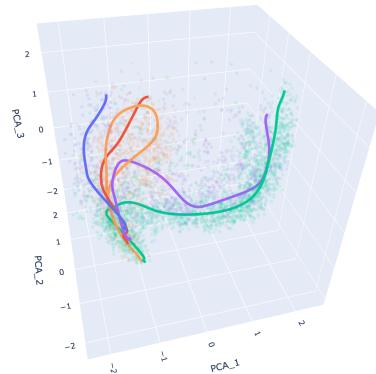
Looking at Third Tranche (11,43,52; Lower left)

Sel Model_Index11

Local Principal Curves for Multiple Perturbations (all)



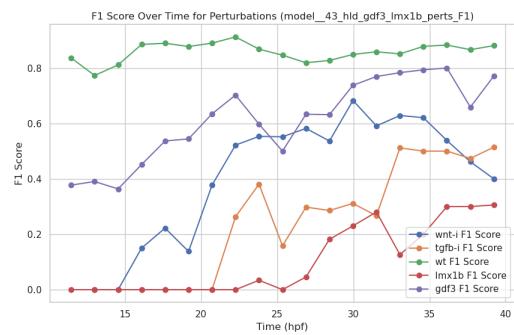
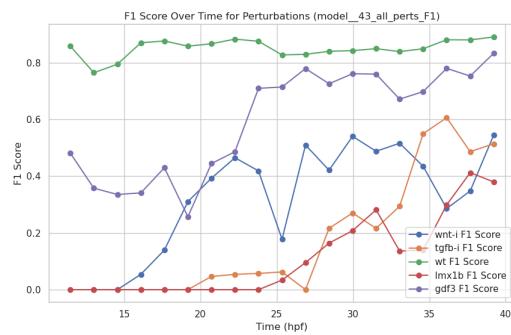
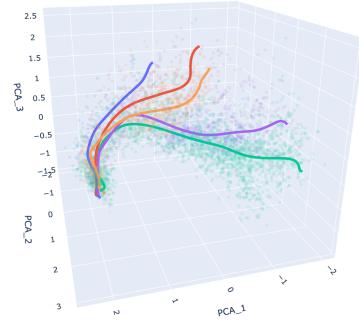
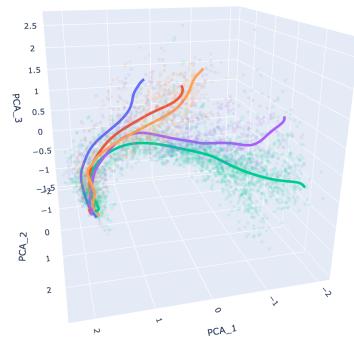
Local Principal Curves for Multiple Perturbations (hld)



Sel Model_Index 43

Local Principal Curves for Multiple Perturbations (all)

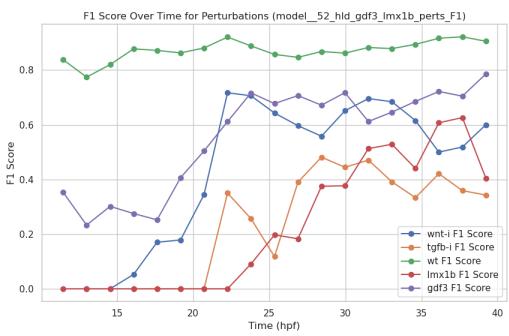
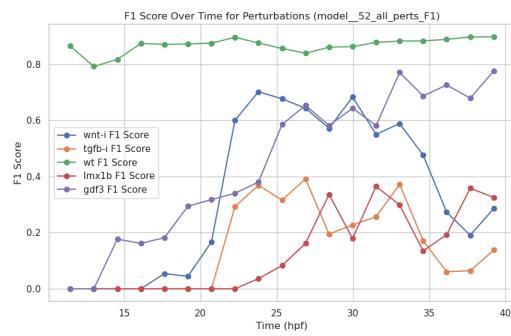
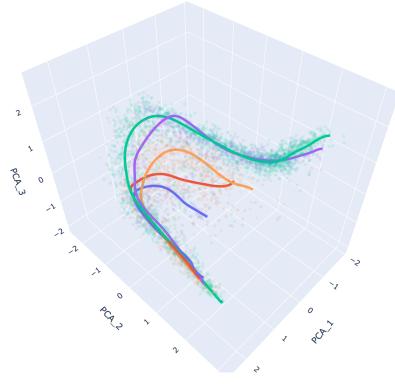
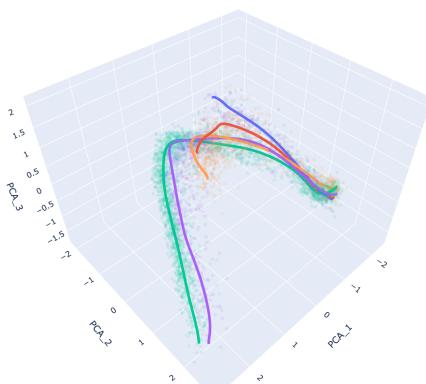
Local Principal Curves for Multiple Perturbations (hid)



Sel Model_Index 52

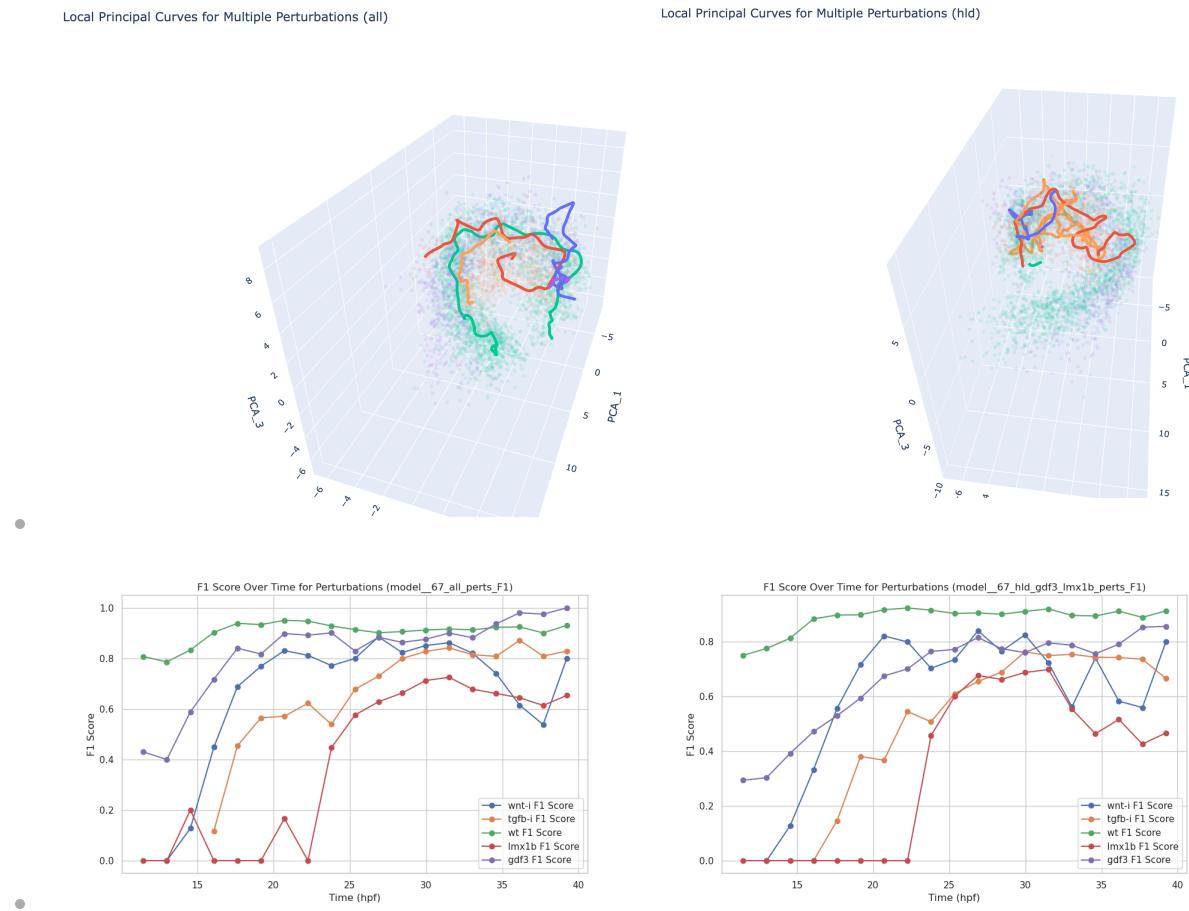
Local Principal Curves for Multiple Perturbations (all)

Local Principal Curves for Multiple Perturbations (hid)

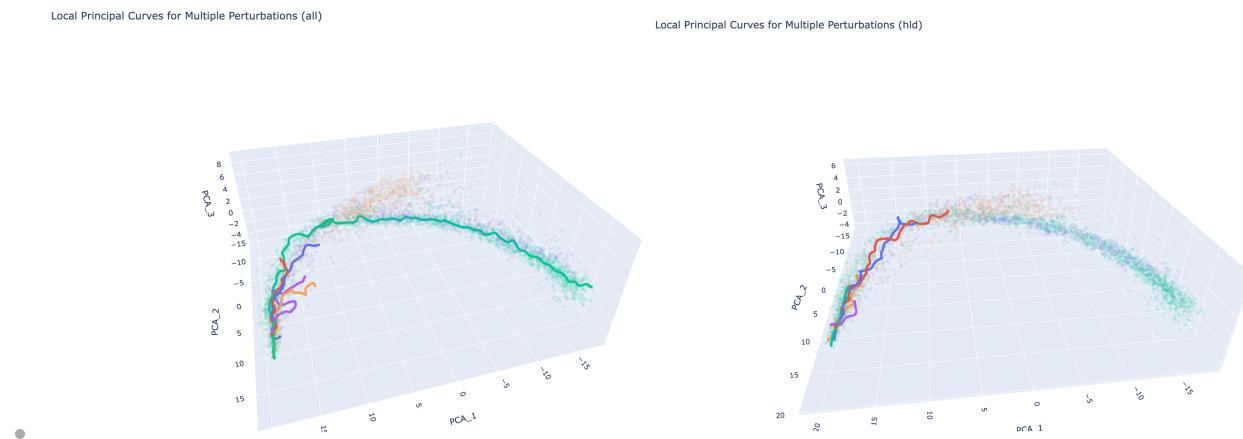


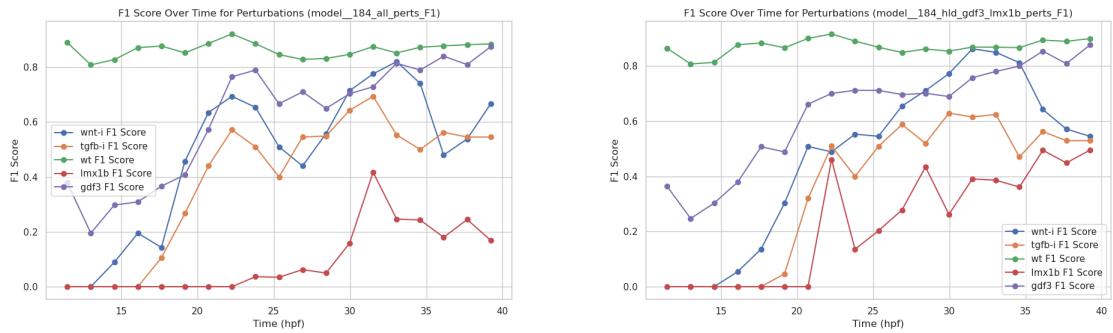
Looking at Fourth Tranche (67, 184, 108; Top Left)

Sel Model_Index 67

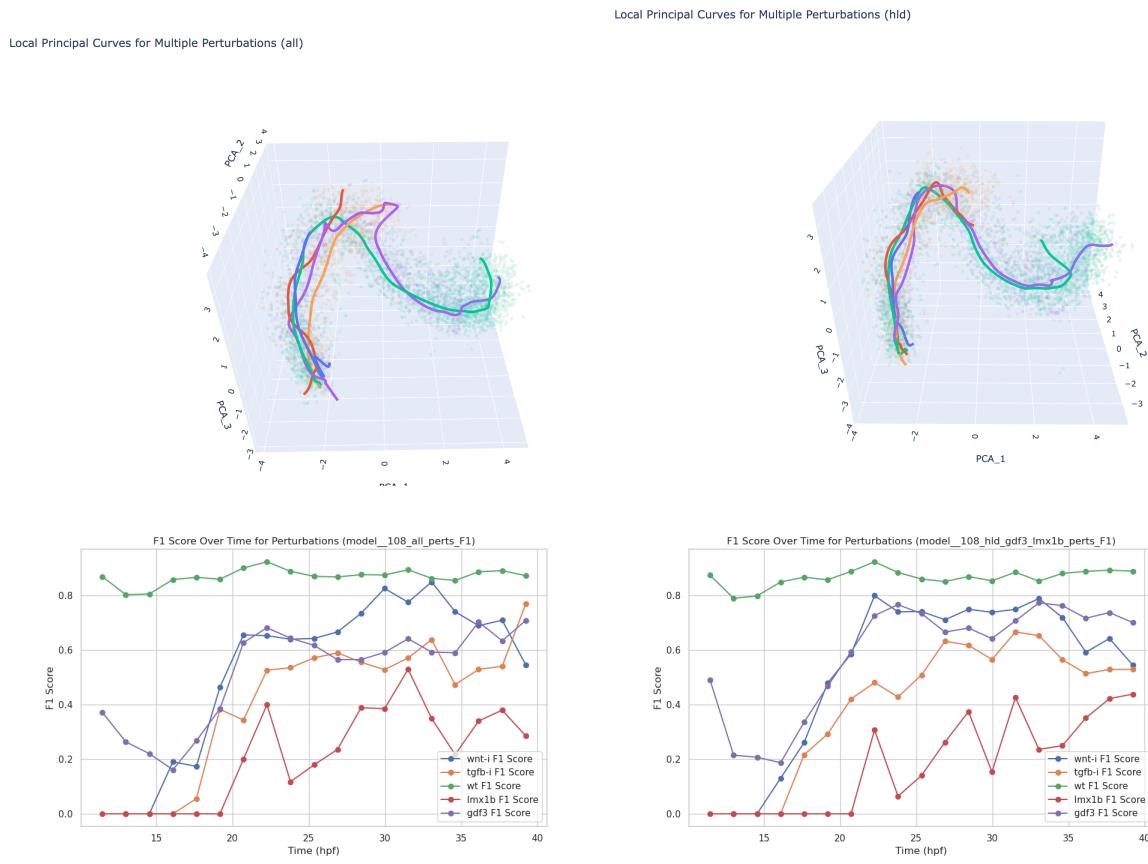


Sel Model_Index 184





Sel Model_Index 108



Choosing Model 74

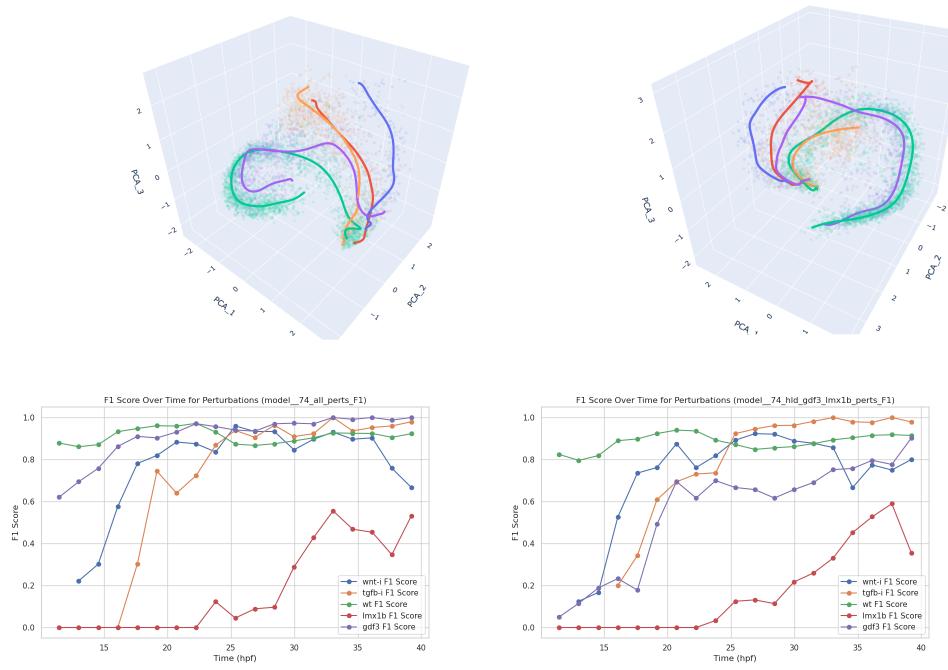
Summary of why model 74

- I choose model 74 because it still has tree like structure, is stable, and it has good classification accuracy
- Reading below we see that it does a better job extracting data than other models e.g. model 97 (not necessarily those in its same tanche e.g. 77)
- It even does better than the previous model we were using!!!

Sel Model_Index 74

Local Principal Curves for Multiple Perturbations (all)

Local Principal Curves for Multiple Perturbations (hld)



comparing it to models with higher classification [MorphSeq](#)
[Param Sweep with Splines; Dec-17-2024 > Looking at Second Traunche \(97,98; Top Right\)](#)

Pretiness:

- Obviously these models with classification just have more compact less and less tree like morphology space (notice how the perts dont branch out)
 - Model 74 wins this easily

Now Classification is more nuanced but it still LOSES:

- The model 97 (and 98) seem to really be edging out the improved F1 scores through improvements in lmx1b.
- However theres two things that should be noted:

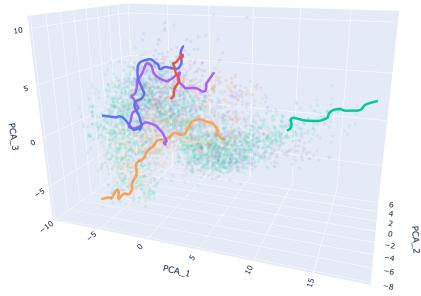
1. That the gains in lmx1b in earlier time points arent learned when held out
2. In all Perts other than lmx1b the model 97 is doing WORSE

- These two observations make me suspicious that it is learning non-biological information. If the model couldve repeated this performance in the hld I wouldve believed it was possibly learning something useful.

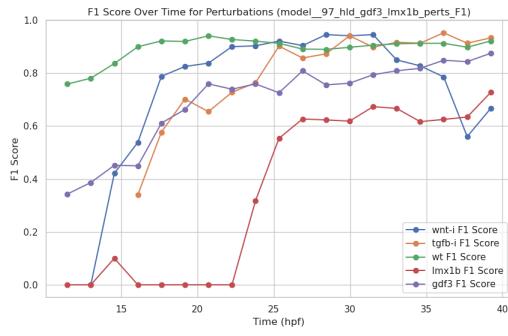
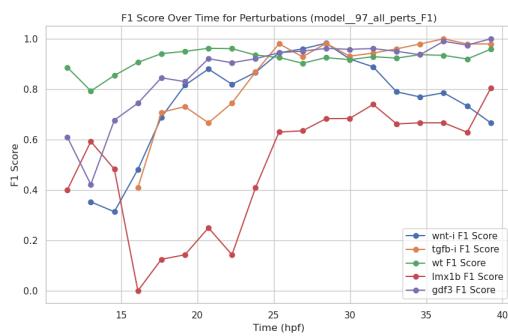
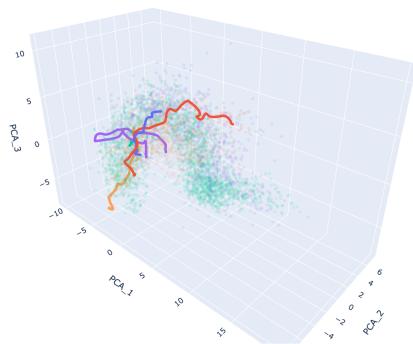
Model 97

Sel Model_Index 97

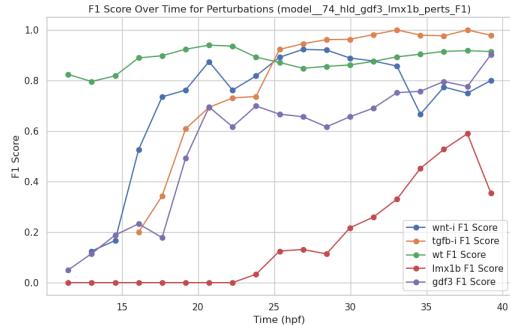
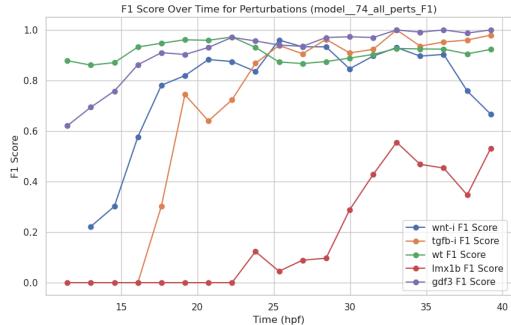
Local Principal Curves for Multiple Perturbations (all)



Local Principal Curves for Multiple Perturbations (hid)

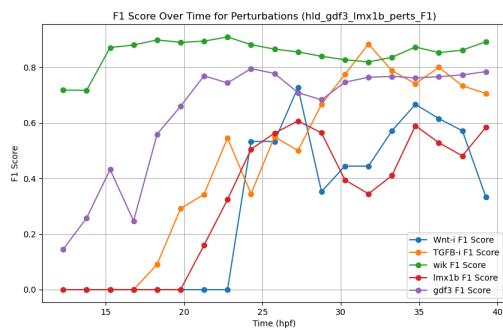
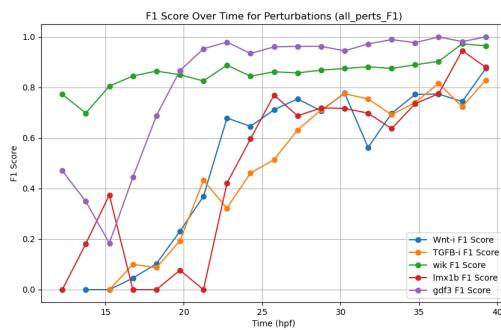


Model 74

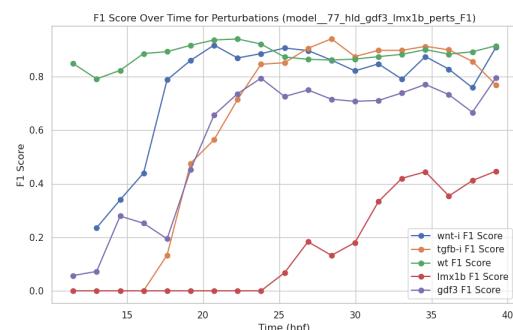
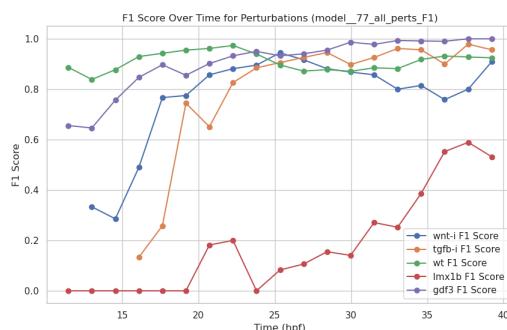


comparing 74 to previous model we were using

- First striking takeaway is that model 74 does a much better job of EXTRACTING relevant information for classification at earlier time points in comparison to previous model. (note that all the best models do a better job not just 74)
 - compare the purple lines (gdf3)
- Prev model:



- Model 74:



Following up on gains on Classification

- I was curious if increasing the tol (tol: 0.05 --> .10e_4) of classification (letting log reg model run longer) will improve results and the answer is not really by much
 - Increasing the tol of classification:

